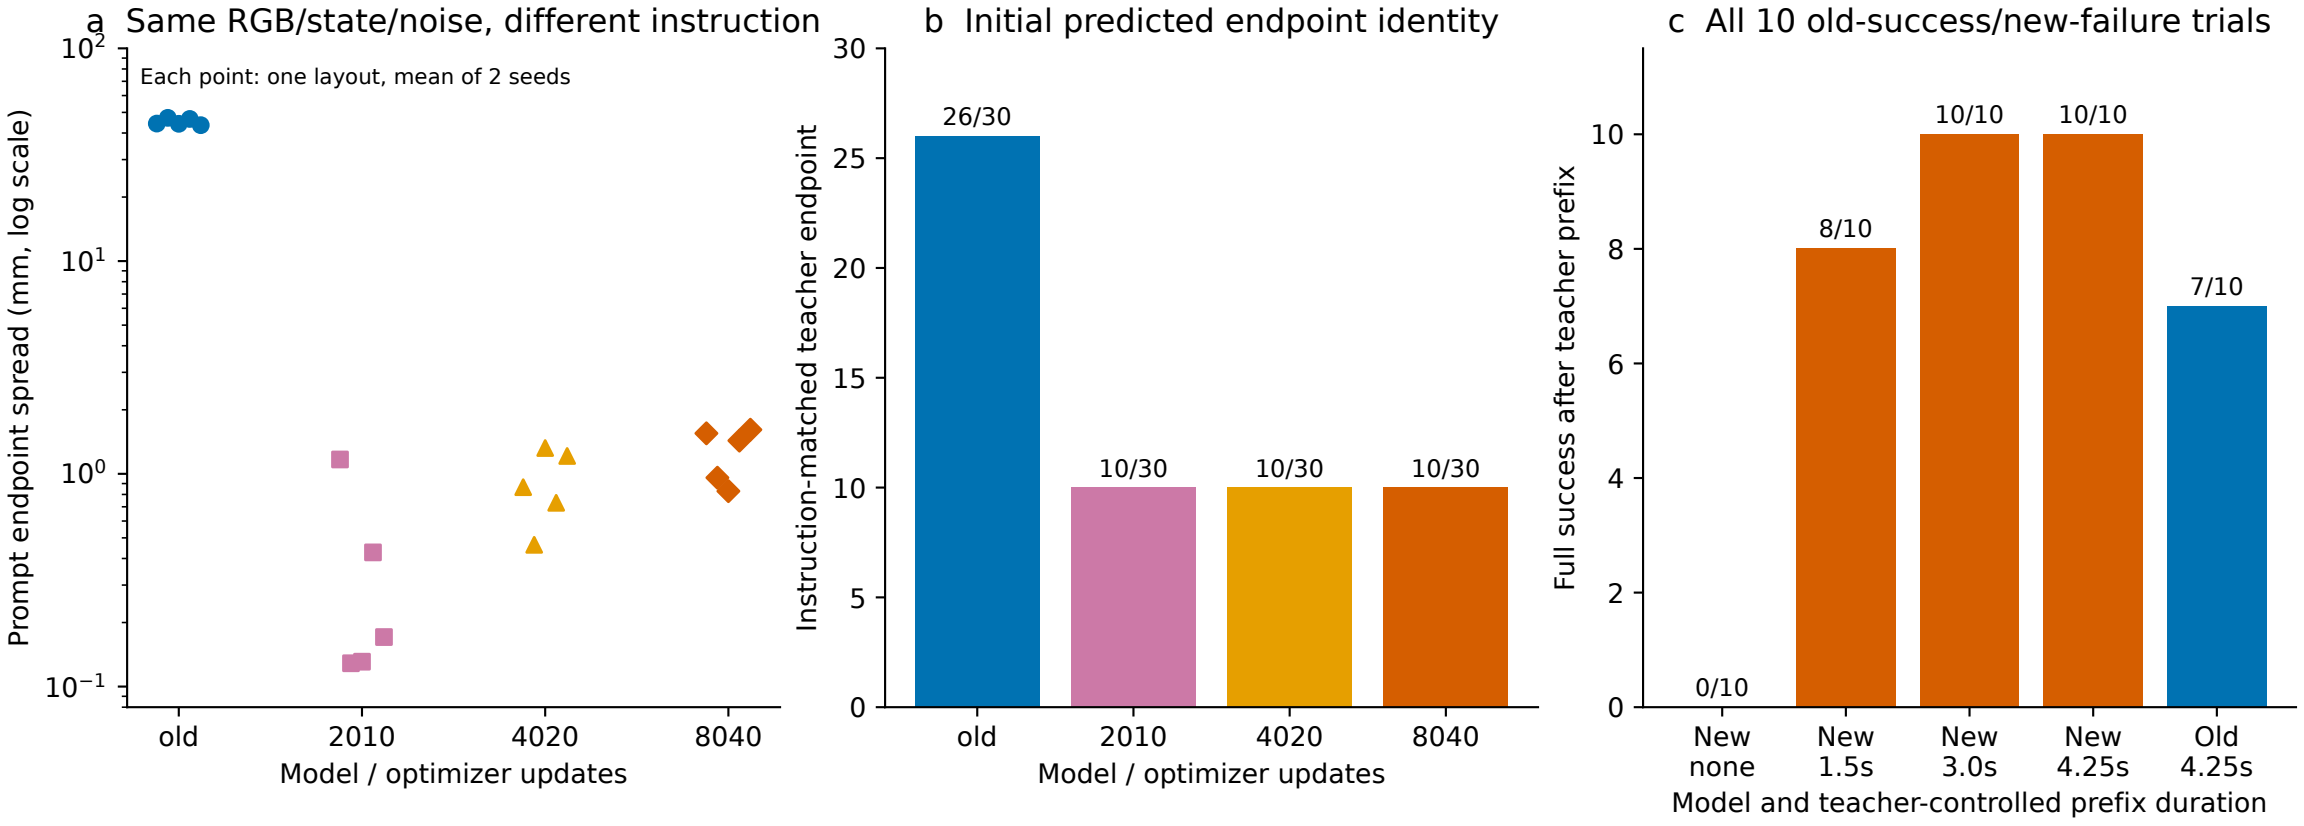
a Same RGB/state/noise, different instruction
Each point: one layout, mean of 2 seeds
Prompt endpoint spread (mm, log scale)
10^2
10^1
10^0
10^{-1}
old
2010
4020
8040
Model / optimizer updates
b Initial predicted endpoint identity
Instruction-matched teacher endpoint
30
25
20
15
10
5
0
26/30
10/30
10/30
10/30
old
2010
4020
8040
Model / optimizer updates
c All 10 old-success/new-failure trials
Full success after teacher prefix
10
8
6
4
2
0
0/10
8/10
10/10
10/10
7/10
New none
New 1.5s
New 3.0s
New 4.25s
Old 4.25s
Model and teacher-controlled prefix duration